

Rohan Vahalia

Washington, DC • (202) 495-1667 • r.vahalia@gmail.com
LinkedIn: RVahalia • GitHub: Riders994

Data engineer and scientist with 7+ years of professional experience and a mathematics degree, drawn to complex systems problems that sit at the intersection of engineering rigor and analytical curiosity. Track record of owning migrations and infrastructure end to end, building for extensibility, mentoring engineers at scale, and partnering closely with the people who depend on the data. A decade-long independent practice building sports data tools rounds out a career built on the belief that good technology gives more people access to a better future.

TECHNICAL SKILLS

Data Warehousing & Modeling: Snowflake (SnowPro Core certified), DBT, dimensional/Kimball modeling, SQL (PostGRES, MySQL)

Orchestration & Infrastructure: Dagster, Airflow, Terraform, Kubernetes, Docker, AWS (S3, EC2)

Programming: Python (pandas, scikit-learn, Keras, Flask), R, SQL, Java

Machine Learning & Statistics: Random Forest, Neural Networks, Clustering, A/B Testing, Elo systems, walk-forward validation

Data Governance: GDPR compliance tooling, data validation and reconciliation, reversible PII anonymization

AI-Assisted Development: Claude Code and Claude — daily use for test generation, edge-case discovery, open-source contribution workflows, and unsupervised-task experimentation (public prompt logs available)

PROFESSIONAL EXPERIENCE

Indeed Inc.

Austin, TX • October 2017 – July 2025

Software Engineer

July 2024 – July 2025

- Extended APIs of TestStats and MCStats, Indeed's internal A/B testing platforms, adding power analysis capabilities and enabling MCStats to delegate calculations to TestStats for improved performance
- Contributed to migration of MCStats' query layer from IQL to Presto for more performant filtering; worked across Python Flask routes and a Java backend

Data Engineer

November 2019 – July 2024

- Led end-to-end data architecture for Indeed's ATS migration from Workday to iCIMS as technical lead on a 3-person team: designed the ERD and bridge-table model merging 2.3 million candidate and 32,000 job records, built custom Python API connectors for a streaming data feed, and engineered an S3-based pagination checkpoint system so large loads could fail and resume gracefully across Kubernetes pods — while extending GDPR erasure tooling to cover every new data relationship
- Drove the POC evaluation and full DAG migration of the Global Revenue Data Mart from Airflow to Dagster: assessed three migration paths, identified a blocker with Indeed's forked Airflow version before it became a costly mistake, worked directly with the Dagster engineering team on undocumented sensor behavior, and designed a cost-aware sensor architecture that minimized Snowflake token usage during parallel validation
- Built and maintained Terraform tooling to automate Snowflake role and permission assignments across a 400+ person engineering org; served as Git admin across 3 teams
- Architected the MSLT dashboard pipeline from its inception: designed Imhotep-vs-Postgres query routing for 40+ global marketing metrics across 10+ countries, built an extensible Python/Google Sheets API tool with documented contribution process, and maintained the system for 5 years
- Rebuilt a fragile, multi-hour manual Jupyter notebook process as a resilient Python tool with built-in retries, structured error reporting, and memory management, eliminating the timeout and data-timing failures that were blocking production use
- Led data development on a sales segmentation tool and data mart enabling Sales Strategy and Client Insights teams to slice a sales representative's book of business, with a \$500M revenue impact
- Conducted ongoing MySQL-to-Postgres and IQL-to-SQL migrations, converted views to materialized views, and maintained dimensional models across the Global Revenue Data Mart as Indeed's product portfolio evolved
- Trained 150+ new BI hires on code review standards (using a critique framework adapted from high-stakes professional fields) and Git practices over a two-year span; conducted 40+ interviews across BI and Data Engineering roles over six years

BI Developer

October 2017 – November 2019

- Built automated Python tooling to calculate and distribute all major marketing metrics required by senior leadership, saving hundreds of person-hours per month
- Developed Python scripts and Jupyter Notebooks to enhance reporting capabilities for marketing and sales teams; participated in early Imhotep-to-Postgres query architecture decisions for what became the MSLT dashboard project

Position2

Santa Clara, CA • April 2015 – February 2017

Data Scientist, Client Manager

- Built repeatable R-based reporting and data-comparison workflows for Fortune 500 clients including FICO and ChooseEnergy; formalized an A/B testing policy integrating Optimizely across all client ad campaigns
- Presented data strategy and Quarterly Business Reports directly to client leadership; pitched analytical services to prospective clients including Malwarebytes and EasyPerk
- Managed Google AdWords bidding strategy, keyword research, and ad copy for major clients; used customer cohort data to evaluate marketing automation and acquisition-funnel performance

EDUCATION

Galvanize Data Science Immersive

Austin, TX • June 2017

American University

Washington, D.C. • May 2014

Bachelor of Science, Mathematics • 2013 William Lowell Putnam Exam score: 10

CERTIFICATIONS

- Snowflake SnowPro Core (COF-C03)
- Cloudera Certified Associate — Apache Spark